

BFS search in miniKanren

KUANG-CHEN LU, Indiana University

WEIXI MA, Indiana University

DANIEL P. FRIEDMAN, Indiana University

The syntax of a programming language should reflect its semantics. When using a disjunction operator in relational programming, a programmer would expect all clauses of this disjunction to share the same chance of being explored, as these clauses are written in parallel. The existing disjunctive operators in miniKanren, however, prioritize their clauses by the order of which these clauses are written down. We have devised a new search strategy that searches evenly in all clauses. Based on our statistics, miniKanren slows down by a constant factor after applying our search strategy. (tested with very-recursiveo, need more tests)

ACM Reference Format:

Kuang-Chen Lu, Weixi Ma, and Daniel P. Friedman. 2019. BFS search in miniKanren. 1, 1 (March 2019), 5 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

miniKanren is a relational programming language embedable in many languages[cite miniKanren.org?].

The version of miniKanren in *The Reasoned Schemer, 2nd Edition* features an efficient and complete search strategy – interleaving depth-first search (iDFS). iDFS biases toward left conde lines. So miniKanren programmers sometimes need to organize their conde lines carefully. We proposed two search strategies and their implementations. The first strategy is breadth-first search. The second one is a modified iDFS.

OUTLINE:

- (About miniKanren)
- (Why the left clauses are explored more frequently?)
- (How to solve the problem?)
- (Summary of later sections)

2 COST OF ANSWERS

The *cost* of an answer is the number of relation applications needed to find the answer. This idea is borrowed from Silvija Seres's work [*]. Now we illustrate the costs of answers by running a miniKanren relation. Fig. 1 defines the relation `repeato` that relates a term `x` with a list whose elements are all `xs`.

Consider the following run of `repeato`.

```
> (run 4 q
    (repeato '* q))
'(( ) (* ) (* *) (* * * ))
```

Authors' addresses: Kuang-Chen LuIndiana University; Weixi MaIndiana University; Daniel P. FriedmanIndiana University.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2019 Association for Computing Machinery.
XXXX-XXXX/2019/3-ART \$15.00
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

```

48 (defrel (repeato x out)
49   (conde
50     [(== '() out)]
51     [(fresh (res)
52       (== '(,x . ,res) out)
53       (repeato x res))]))
54
55
56

```

Fig. 1. `repeato`

The above `run` generates 4 answers. All are lists of `*s`. The order of the answers reflects the order miniKanren discovers them: the leftmost answer is the first one. This result is not surprising: to generate the first answer, `'()`, miniKanren needs to apply `repeato` only once and the later answers need more recursive applications. In this example, the cost of each answer is the same as one more than the number of `*s`: the cost of `'()` is 1, the cost of `'(*)` is 2, and so on.

A list of answer is in the *cost-respecting* order if no answer occurs before another answer of a lower cost. In the above example, the answers are cost-respecting. The iDFS search, however, does not generate cost-respecting answers in general. As an example, consider the following `run` of `repeato`.

```

67 > (run 12 q
68   (conde
69     [(repeato 'a q)]
70     [(repeato 'b q)]
71     [(repeato 'c q)]))
72 '(() (a) ()
73   (a a) () (a a a)
74   (b) (a a a a) (c)
75   (a a a a a) (b b) (a a a a a a))
76

```

The results are not cost-respecting. For example, `'(a a)` occurs before `'(b)` while `'(a a)` is associated with a higher cost. The problem is that iDFS strategy prioritizes the first `conde` case considerably. In general, when every `conde` case are equally productive, the iDFS strategy takes $1/2^i$ answers from the i -th case, except the last case, which share the same portion as the second last.

For the above `run`, both search strategies produces answers in increasing order of costs, i.e. both of them are *cost-respecting*. In more complicated cases, however, interleaving DFS might not produce answers in cost-respecting order. For instance, with iDFS the `run` in Fig. ?? produces answers in a seemingly random order. In contrast, the same run with BFS produces answers in an expected order (Fig. ??).

```

85 > (run 12 q
86   (conde
87     [(repeato 'a q)]
88     [(repeato 'b q)]
89     [(repeato 'c q)]))
90 '(() () ()
91   (a) (b) (c)
92   (a a) (b b) (c c)
93
94

```

```

95 (define (append-inf s-inf t-inf)
96   (cond
97     ((null? s-inf) t-inf)
98     ((pair? s-inf)
99      (cons (car s-inf)
100            (append-inf (cdr s-inf) t-inf)))
101     (else (lambda ()
102              (append-inf t-inf (s-inf))))))
103
104

```

Fig. 2. `append-inf` in `mk-0`

```

108 (a a a) (b b b) (c c c)
109
110

```

3 CHANGE SEARCH STRATEGY

Now we change the search strategy and optimize the system. The whole process is completed in three steps, corresponding to 4 versions of miniKanren. The initial version, `mk-0`, is exactly the miniKanren in *The Reasoned Schemer, 2nd Edition*.

3.1 from `mk-0` to `mk-1`

In `mk-0` and `mk-1`, search spaces are represented by streams of answers. Streams can be finite or infinite. Finite streams are just lists. And infinite streams are improper lists, whose last `cdr` is a thunk returning another stream. We call the `cars` the *mature* part, and the last `cdr` the *immature* part.

Streams are cost respective when they are initially constructed by `==`. However, the `mk-0` version of `append-inf` (Fig. 2) breaks cost respectiveness when its first argument, `s-inf`, is infinite. The resulting mature part contains only the mature part of `s-inf`. The whole `t-inf` goes to the resulting immature part.

The `mk-1` version of `append-inf` (Fig. 3) restores cost-respectiveness by combining the mature parts in the fashion of `append`. This `append-inf` calls its helper immediately, with the first argument, `s?`, set to `#t`, which means `s-inf` in the helper is the `s-inf` in the driver. Two streams are swapped in the third `cond` clause, where `s?` is also changed accordingly.

`mk-1` is not efficient in two aspects. `append-inf` need to copy all `cons` cells of two input stream when the first one is infinite. Besides, `mk-1` generates answers of the same cost at once, even when only a small portion is queried. We solves the two problems in the next two subsections.

3.2 stepstone of optimization

In `mk-2` we set up an interface between streams and the rest of miniKanren, so that the optimization in `mk-3` is easier to present. As a side-effect, we also solve the problem of copying too many `cons` cells. We list all interface functions in Fig. 4.

The representation of stream in `mk-2` is a pair of mature part and immature part. The mature parts are lists of answers. And the immature part is either a thunk or `'()`.

Changes in `mk-2` are mostly about setting up the interface, except the change in `append-inf` (Fig. 5). This version of `mk-2` essentially does the same thing as the `mk-1` version.

```

142 (define (append-inf s-inf t-inf)
143   (append-inf^ #t s-inf t-inf))
144
145 (define (append-inf^ s? s-inf t-inf)
146   (cond
147     ((pair? s-inf)
148      (cons (car s-inf)
149            (append-inf^ s? (cdr s-inf) t-inf)))
150     ((null? s-inf) t-inf)
151     (s? (append-inf^ #f t-inf s-inf))
152     (else (lambda ()
153              (append-inf (t-inf) (s-inf))))))
154
155
156

```

Fig. 3. append-inf in mk-1

interface function	meaning
empty-inf	construct an empty stream
unit-mature-inf	construct a stream with one answer
unit-immature-stream	construct a stream with one thunk
null-inf?	predicate whether a stream is empty
mature-inf?	predicate whether a stream has some mature answers
car-inf	take the first answer out of a mature stream
cdr-inf	remove the first answer of a mature stream
force-inf	get answers of higher cost

Fig. 4. interface functions and their meanings in mk-2 and mk-3

```

172 (define (append-inf s-inf t-inf)
173   (cons (append (car s-inf) (car t-inf))
174         (let ((s (cdr s-inf))
175               (t (cdr t-inf)))
176           (cond
177             ((null? s) t)
178             ((null? t) s)
179             (else (lambda () (append-inf (s) (t))))))))
180
181
182

```

Fig. 5. append-inf in mk-2

3.3 optimization

The goal is to express BFS explicitly with queue, so that the system doesn't generate all answers of the same cost at once.

Interesting changes: (1) put thunks in a list; (2) change force-inf (introduced in 4.B) so that it can make progress in all thunks (3) use a queue to manage thunks in take-inf.

4 CONCLUSION

ACKNOWLEDGMENTS

REFERENCES