miniKanren with fair search strategies

KUANG-CHEN LU, Indiana University

WEIXI MA, Indiana University

DANIEL P. FRIEDMAN, Indiana University

The syntax of a programming language should reflect its semantics. When writing a $cond^e$ expression in miniKanren, a programmer would expect all clauses share the same chance of being explored, as these clauses are written in parallel. The existing search strategy, interleaving depth-first search (DFS_i), however, prioritize its clauses by the order how they are written down. Similarly, when a $cond^e$ is followed by another goal conjunctively, a programmer would expect answers in parallel share the same chance of being explored. Again, the answers by DFS_i is different from the expectation. We have devised three new search strategies that have different level of fairness in `disj` and `conj`.
   c1

1 INTRODUCTION

miniKanren is a family of relational programming languages. Friedman et al. [2] introduce miniKanren and its implementation in *The Reasoned Schemer, 2nd Ed* (TR2). miniKanren programs, especially relational interpreters, have been proven to be useful in solving many problems by Byrd et al. [1].

A subtlety arises when a $cond^e$ contains many clauses: not every clause has an equal chance to contribute to the result. As an example, consider the following relation `repeato` and its invocation.

```
(defrel (repeato x out)
   (conde
     [(== `(,x) out)]
     [(fresh (res)
        (== `(,x . ,res) out)
        (repeato x res))]))
> (run 4 q
    (repeato '* q))
'((*) (* *) (* * *) (* * * *))
```

Next, consider the following disjunction of invoking `repeato` with four different letters.

[c1] *LKC*: `disj` and `conj` occurs free in the abstract.

Authors' addresses: Kuang-Chen LuIndiana University; Weixi MaIndiana University; Daniel P. FriedmanIndiana University.

Unpublished working draft. Not for distribution.

```
> (run 12 q
    (conde
      [(repeato 'a q)]
      [(repeato 'b q)]
      [(repeato 'c q)]
      [(repeato 'd q)]))
```

$cond^e$ intuitively relates its clauses with logical or. And thus an unsuspicious beginner would expect each letter to contribute equally to the result, as follows.

```
'((a) (b) (c) (d)
  (a a) (b b) (c c) (d d)
  (a a a) (b b b) (c c c) (d d d))
```

The $cond^e$ in TR2, however, generates a less expected result.

```
'((a) (a a) (b) (a a a)
  (a a a a) (b b)
  (a a a a a) (c)
  (a a a a a a) (b b b)
  (a a a a a a a) (d))
```

The miniKanren in TR2 implements interleaving DFS (DFS_i), the cause of this unexpected result. With this search strategy, each clause takes half of its received computational resources and pass the other half to its following clauses, except for the last clause that takes all resources it receives. In the example above, the a clause takes half of all recourses. And the b clause takes a quarter. Thus c and d barely contribute to the result.

DFS_i is sometimes powerful for an expert. By carefully organizing the order of $cond^e$ clauses, a miniKanren program can explore more "interesting" clauses than those uninteresting ones, and thus use computational resources efficiently. A little miniKanrener, however, may beg to differ–understanding implementation details and fiddling with clauses order is not the first priority of a beginner.

There is another reason that miniKanren could use more search strategies than just DFS_i. In many applications, there does not exist one order that serves for all purposes. For example, a relational dependent type checker contains clauses for constructors that build data and clauses for eliminators that use data. When the type checker is used to generate simple and shallow programs, the clauses of constructors should be put in the front of $cond^e$. When performing proof searches for complicated programs, the clauses of eliminator should take the focus. With DFS_i, these two uses cannot be efficient at the same time. In fact, to make one use efficient, the other one must be drastically slow.

The specification that every clause in the same $cond^e$ is given equal "search priority" is called fair disj. And search strategies with almost-fair disj give every clause in the same $cond^e$ similar priority. Fair conj, a related concept, is more complicated. We defer it to the next section.

To summarize our contribution, we

- propose and implement balanced interleaving depth-first search (DFS_{bi}), a new search strategy with almost-fair disj.
- propose and implement fair depth-first search (DFS_f), a new search strategy with fair disj.
- implement in a new way breath-first search (BFS), a search strategy with fair disj and fair conj. (our code runs faster in all benchmarks and is simpler). And we prove formally that our BFS implementation is equivalent to the one by Seres et al. [5].

2 SEARCH STRATEGIES AND FAIRNESS

In this section, we define fair disj, almost-fair disj and fair conj. Before going further into fairness, we would like to give a short review about state, search space, and goal, because fairness is defined in terms of them. A *state* is a collection of constraints. Every answer corresponds to a state. A *search space* is a collection of states. And a *goal* is a function from a state to a search space.

Now we elaborate fairness by running more queries about repeato.

2.1 fair disj

Given the following program, it is natural to expect lists of each letter to constitute 1/4 in the answer. DFS_i, the current search strategy, however, results in many more lists of as than lists of other letters. And some letters (e.g. c and d) are rarely seen. The situation would be exacerbated if cond^e contains more clauses.

```
;; DFSi (unfair disj)
> (run 12 q
    (conde
      ((repeato 'a q))
      ((repeato 'b q))
      ((repeato 'c q))
      ((repeato 'd q))))
'((a) (a a) (b) (a a a)
  (a a a a) (b b)
  (a a a a a) (c)
  (a a a a a a) (b b b)
  (a a a a a a) (d))
```

Under the hood, the cond^e here is allocating computational resource to four trivially different search space. The unfair disj in DFS_i allocates many more resources to the first search space. On the contrary, fair disj would allocate resources evenly to each search space.

```
;; DFSf (fair disj)
> (run 12 q
    (conde
      ((repeato 'a q))
      ((repeato 'b q))
      ((repeato 'c q))
      ((repeato 'd q))))
'((a) (b) (c) (d)
  (a a) (b b) (c c) (d d)
  (a a a) (b b b) (c c c) (d d d))
```

```
                        ;; BFS (fair disj)
                      > (run 12 q
                          (conde
                            ((repeato 'a q))
                            ((repeato 'b q))
                            ((repeato 'c q))
                            ((repeato 'd q))))
                        '((a) (b) (c) (d)
                          (a a) (b b) (c c) (d d)
                          (a a a) (b b b) (c c c) (d d d))
```

Running the same program again with almost-fair disj (e.g. DFS$_{bi}$) gives the same result. Almost-fair, however, is not completely fair, as shown by the following example.

```
                        ;; biDFS (almost-fair disj)
                      > (run 16 q
                          (conde
                            [(repeato 'a q)]
                            [(repeato 'b q)]
                            [(repeato 'c q)]
                            [(repeato 'd q)]
                            [(repeato 'e q)]))
                        '((b) (c) (d) (a)
                          (b b) (c c) (d d) (e)
                          (b b b) (c c c) (d d d) (a a)
                          (b b b b) (c c c c) (d d d d) (e e))
```

DFS$_{bi}$ is fair only when the number of goals is a power of 2, otherwise, some goals would be allocated twice as many resources than the others. In the above example, where the cond^e has five clauses, the clauses of b, c, and d are allocated more resources.

We end this subsection with precise definitions of all levels of disj fairness. Our definition of *fair* disj is slightly generalize from the one by Seres et al. [5]. Their definition is for binary disjunction. We generalize it to multi-arity one.

DEFINITION 2.1 (FAIR disj). *A disj is fair if and only if it allocates computational resource evenly to search spaces produced by goals in the same disjunction (i.e. clauses in the same* cond^e *).*

DEFINITION 2.2 (ALMOST-FAIR disj). *A disj is almost-fair if and only if it allocates computational resource so evenly to search spaces produced by goals in the same disjunction that the maximal ratio of resources is bounded by a constant.*

DEFINITION 2.3 (UNFAIR disj). *A disj is unfair if and only if it is not even almost-fair.*

2.2 fair conj

c1

---

c1 *LKC*: Checked grammer of text before this line

Given the following program, it is natural to expect lists of each letter to constitute 1/4 in the answer. Search strategies with unfair $conj$ (e.g. DFS_i, DFS_{bi}, DFS_f), however, results in many more lists of as than lists of other letters. And some letters are rarely seen. The situation would be exacerbated if $cond^e$ contains more clauses.

[c2]

```
;; DFSi (unfair conj)
> (run 12 q
    (fresh (x)
      (conde
        ((== 'a x))
        ((== 'b x))
        ((== 'c x))
        ((== 'd x)))
      (repeato x q)))
'((a) (a a) (b) (a a a)
  (a a a a) (b b)
  (a a a a a) (c)
  (a a a a a a) (b b b)
  (a a a a a a a) (d))


;; DFSf (unfair conj)
> (run 12 q
    (fresh (x)
      (conde
        ((== 'a x))
        ((== 'b x))
        ((== 'c x))
        ((== 'd x)))
      (repeato x q)))
'((a) (a a) (b) (a a a)
  (a a a a) (b b)
  (a a a a a) (c)
  (a a a a a a) (b b b)
  (a a a a a a a) (d))
```

[c2] *LKC*: Should I add a footnote?

" Although DFS_i's $disj$ is unfair in general, it is fair when there is no call to relational definition in $cond^e$ clauses (including this case). "

```
;; DFSbi (unfair conj)
> (run 12 q
    (fresh (x)
      (conde
        ((== 'a x))
        ((== 'b x))
        ((== 'c x))
        ((== 'd x)))
      (repeato x q)))
'((a) (a a) (c) (a a a)
  (a a a a) (c c)
  (a a a a a) (b)
  (a a a a a a) (c c c)
  (a a a a a a a) (d))
```

Under the hood, the cond^e and the call to repeato are connected by conj. The cond^e goal outputs a search space including four trivially different states. Then the next conjunctive goal, (repeato x q), is applied to each of these states, producing four trivially different search spaces. In the examples above, the conjs are allocating more computational resources to the search space of a. On the contrary, fair conj would allocate resources evenly to each search space.

```
;; BFS (fair conj)
> (run 12 q
    (fresh (x)
      (conde
        [(== 'a x)]
        [(== 'b x)]
        [(== 'c x)]
        [(== 'd x)])
      (repeato x q)))
'((a) (b) (c) (d)
  (a a) (b b) (c c) (d d)
  (a a a) (b b b) (c c c) (d d d))
```

A more interesting situation is when the first conjunct produces infinite many answers. Consider the following example, a naive specification of fair conj might require search strategies to produce all sorts of singleton lists, but no longer ones, which makes the strategies incomplete. A search strategy is *complete* if and only if it can find out all the answers within finite time, otherwise, it is *incomplete*.

```
;; naively fair conj
> (run 6 q
    (fresh (xs)
      (conde
        [(repeato 'a xs)]
        [(repeato 'b xs)])
      (repeato xs q)))
'(((a)) ((b))
  ((a a)) ((b b))
  ((a a a)) ((b b b)))
```

Our solution requires a search strategy with *fair* conj to organize states in bags in search spaces, where each bag contains finite states, and to allocate resources evenly among search spaces derived from states in the same bag. It is up to a search strategy designer to decide by what criteria to put states in the same bag, and how to allocate resources among search spaces related to different bags.

BFS puts states of the same cost in the same bag, and allocate resources carefully among search spaces related to different bags such that answers are produced in increasing order of cost. The *cost* of a answer is its depth in the search tree (i.e. the number of calls to relational definitions required to find them) [5]. In the following example, the cost of each answer is equal to the length of the inner lists plus the length of the outer list.

```
;; BFS (fair conj)
> (run 12 q
    (fresh (xs)
      (conde
        [(repeato 'a xs)]
        [(repeato 'b xs)])
      (repeato xs q)))
'(((a)) ((b))
  ((a) (a)) ((b) (b))
  ((a a)) ((b b))
  ((a) (a) (a)) ((b) (b) (b))
  ((a a) (a a)) ((b b) (b b))
  ((a a a)) ((b b b)))
```

We end this subsection with precise definitions of all levels of conj fairness.

DEFINITION 2.4 (FAIR conj). *A* conj *is fair if and only if it allocates computational resource evenly to search spaces produced from states in the same bag. A bag is a finite collection of state. And search strategies with fair* conj *should represent search spaces with possibly infinite collection of state.*

DEFINITION 2.5 (UNFAIR conj). *A* conj *is unfair if and only if it is not fair.*

3 INTERLEAVING DEPTH-FIRST SEARCH

In this section, we review the implementation of interleaving depth-first search (DFS$_i$). This review is for comparison with other search strategies in this paper. Thus we focus on parts that are changed later. TRS2 [2] provides a comprehensive description of the whole miniKanren implementation.

```
330  #| Goal x Goal -> Goal |#
331  (define (disj2 g1 g2)
332    (lambda (s)
333      (append-inf (g1 s) (g2 s))))
334
335
336  #| Space x Space -> Space |#
337  (define (append-inf s-inf t-inf)
338    (cond
339      ((null? s-inf) t-inf)
340      ((pair? s-inf)
341       (cons (car s-inf)
342         (append-inf (cdr s-inf) t-inf)))
343      (else (lambda ()
344             (append-inf t-inf (s-inf))))))
345
346
347  #| Goal x Goal -> Goal |#
348  (define (conj2 g1 g2)
349    (lambda (s)
350      (append-map-inf g2 (g1 s))))
351
352  #| Goal x Space -> Space |#
353  (define (append-map-inf g s-inf)
354    (cond
355      ((null? s-inf) '())
356      ((pair? s-inf)
357       (append-inf (g (car s-inf))
358         (append-map-inf g (cdr s-inf))))
359      (else (lambda ()
360             (append-map-inf g (s-inf))))))
361
362
363  (define-syntax disj
364    (syntax-rules ()
365      ((disj) fail)
366      ((disj g) g)
367      ((disj g0 g ...) (disj2 g0 (disj g ...)))))
368
369  (define-syntax conj
370    (syntax-rules ()
371      ((conj) succeed)
372      ((conj g) g)
373      ((conj g0 g ...) (conj2 g0 (conj g ...)))))
374
375
376  (define-syntax conde
       (syntax-rules ()
         ((conde (g ...) ...)
          (disj (conj g ...) ...))))
```

Fig. 1. implementation of DFS$_i$

```
377  (define-syntax disj
378    (syntax-rules ()
379      [(disj) fail]
380      [(disj g ...) (disj+ (g ...) () ())]))
381
382
383  (define-syntax disj+
384    (syntax-rules ()
385      [(disj+ (g) () ()) g]
386      [(disj+ () (gl ...) (gr ...))
387       (disj2 (disj+ (gl ...) () ())
388              (disj+ (gr ...) () ()))]
389      [(disj+ (g0) (gl ...) (gr ...))
390       (disj2 (disj+ (gl ... g0) () ())
391              (disj+ (gr ...) () ()))]
392      [(disj+ (g0 g1 g ...) (gl ...) (gr ...))
393       (disj+ (g ...) (gl ... g0) (gr ... g1))]))
394
```

Fig. 2. implementation of DFS$_{bi}$

Fig. 1 shows part of the implementation of DFS$_i$. We follow a convention to name variables bound to states with 's', to name variables bound to goals with 'g', and to name variables bound to search spaces with a suffix '-inf'. The first function, disj2, implements binary disjunction. append-inf is its helper, which composes two disjunctive search spaces. The following function, conj2, implements binary conjunction. It applies the *first* goal to the input state, then applies the second goal to states in the resulting search space. The latter process is done with a helper function. append-map-inf applies its input goal to states in its input search spaces and compose the resulting search spaces. It reuses append-inf for search space composition. The following three definitions introduce syntactic sugars that miniKanren users are more familiar with. The first two definitions say disjunction and disjunction are right-associative. The next and last definition says cond^e relates its clauses disjunctively, and goals in the same clause conjunctively.

4 BALANCED INTERLEAVING DEPTH-FIRST SEARCH

Balanced interleaving DFS (DFS$_{bi}$) has almost-fair disj and unfair conj. The implementation of DFS$_{bi}$ differs from DFS$_i$'s in the disj macro. We list the new disj with its helper in Fig. 2. The helper disj+ builds a balanced binary tree whose leaves are the goals and whose nodes are disj2s, hence the name of this search strategy. The first argument to disj+ is a list of all goals. And the next two arguments accumulate goals in the left and right subtrees. The first clause says that when there is one goal, the tree is the goal itself. When there are more goals, the first argument is partitioned into two sub-lists. The partition is done by repetitively dispatching the first two goals (the last clause), until no goal remains (the second clause) or one goal remains (the third clause). In contrast, the disj in DFS$_i$ constructs the binary tree with the same collection of leaf nodes but in a particularly unbalanced form.

```
424  #| Goal x Goal -> Goal |#
425  (define (disj2 g1 g2)
426    (lambda (s)
427      (append-inf/fair (g1 s) (g2 s))))
428
429
430  #| Space x Space -> Space |#
431  (define (append-inf/fair s-inf t-inf)
432    (let loop ([s? #t]
433               [s-inf s-inf]
434               [t-inf t-inf])
435      (cond
436        ((pair? s-inf)
437         (cons (car s-inf)
438           (loop s? (cdr s-inf) t-inf)))
439        ((null? s-inf) t-inf)
440        (s? (loop #f t-inf s-inf))
441        (else (lambda ()
442                (append-inf/fair (t-inf) (s-inf)))))))
443
444
445
```

Fig. 3. implementation of DFS$_f$

5 FAIR DEPTH-FIRST SEARCH

Fair DFS (DFS$_f$) has fair `disj` and unfair `conj`. The implementation of DFS$_f$ differs from DFS$_i$'s in `disj2` (Fig. 3). `disj2` is changed to call a new and fair version of `append-inf`. `append-inf/fair` immediately calls its helper, `loop`, with the first argument, `s?`, set to `#t`, which indicates that `s-inf` and `t-inf` haven't been swapped. The swapping happens at the third cond clause in the helper, where `s?` is updated accordingly. The first two cond clauses essentially copy the `car`s and stop recursion when one of the input spaces is obviously finite. The third clause, as we mentioned above, is just for swapping. When the fourth and last clause runs, we know that both `s-inf` and `t-inf` are ended with a thunk. In this case, a new thunk is constructed. The new thunk calls the driver recursively. Here changing the order of `t-inf` and `s-inf` won't hurt the fairness (though it will change the order of answers). We swap them back so that answers are produced in a more natural order.

6 BREADTH-FIRST SEARCH

BFS has both fair `disj` and fair `conj`. Our implementation is based on DFS$_f$ (not DFS$_i$). To implement BFS based on DFS$_f$, we need `append-map-inf/fair` in addition to `append-inf/fair`. The only difference between `append-map-inf/fair` and `append-map-inf` is that the former calls `append-inf/fair` instead of `append-inf`.

The implementation can be improved in two ways. First, as mentioned in section 2.2, BFS puts answers in bags and answers of the same cost are in the same bag. In this implementation, however, it is unclear where this information is recorded. Second, `append-inf/fair` is extravagant in memory usage. It makes $O(n + m)$ new cons cells every time, where n and m are the "length"s of input search spaces. We address these issues in the first subsection.

Both our BFS and Seres's BFS Seres et al. [5] produce answers in increasing order of cost. So it is interesting to see if they are equivalent. We prove so in Coq. The details are in the second subsection.

```
471  (define (none)   '(()   . #f))
472  (define (unit s) '((,s) . #f))
473  (define (step f) '(()   . ,f))
474
475
476  #| Space x Space -> Space |#
477  (define (append-inf/fair s-inf t-inf)
478    (cons (append (car s-inf) (car t-inf))
479      (let ([t1 (cdr s-inf)]
480            [t2 (cdr t-inf)])
481        (cond
482          [(not t1) t2]
483          [(not t2) t1]
484          [else (lambda () (append-inf/fair (t1) (t2)))]))))
485
486  #| Goal x Space -> Space |#
487  (define (append-map-inf/fair g s-inf)
488    (foldr
489      (lambda (s t-inf)
490        (append-inf/fair (g s) t-inf))
491      (let ([f (cdr s-inf)])
492        (step (and f (lambda () (append-map-inf/fair g (f))))))
493      (car s-inf)))
494
495
496  #| option Nat x Space -> [State] |#
497  (define (take-inf n s-inf)
498    (let loop ([n n]
499              [vs (car s-inf)])
500      (cond
501        ((and n (zero? n)) '())
502        ((pair? vs)
503         (cons (car vs)
504           (loop (and n (sub1 n)) (cdr vs))))
505        (else
506         (let ([f (cdr s-inf)])
507           (if f (take-inf n (f)) '()))))))
```

Fig. 4. new and changed functions in optimized BFS that implements pure features

6.1 optimized BFS

As mentioned in section 2.2, BFS puts answers in bags and answers of the same cost are in the same bag. The cost information is recorded subtly – the cars of a search space have cost 0 (i.e. they are in the same bag), and the costs of answers in thunk are computed recursively then increased by one. It is even more subtle that append-inf/fair

```scheme
(define (elim s-inf ks kf)
  (let ([ss (car s-inf)]
        [f (cdr s-inf)])
    (cond
      [(and (null? ss) f)
       (step (lambda () (elim (f) ks kf)))]
      [(null? ss) (kf)]
      [else (ks (car ss) (cons (cdr ss) f))])))

(define (ifte g1 g2 g3)
  (lambda (s)
    (elim (g1 s)
      (lambda (s0 s-inf)
        (append-map-inf/fair g2
          (append-inf/fair (unit s0) s-inf)))
      (lambda () (g3 s)))))

(define (once g)
  (lambda (s)
    (elim (g s)
      (lambda (s0 s-inf) (unit s0))
      (lambda () (none)))))
```

Fig. 5. new and changed functions in optimized BFS that implements impure features

and the append-map-inf/fair respects the cost information. We make these facts more obvious by changing the type of search space, modifying related function definitions, and introducing a few more functions.

The new type is a pair whose car is a list of answers (the bag), and whose cdr is either a #f or a thunk returning a search space. A falsy cdr means the search space is obviously finite.

Functions related to the pure subset are listed in Fig. 4 (the others in Fig. 5). They are compared with Seres et al.'s implementation later. The first three functions in Fig. 4 are search space constructors. none makes an empty search space; unit makes a space from one answer; and step makes a space from a thunk. The remaining functions do the same thing as before.

Luckily, the change in append-inf/fair also fixes the miserable space extravagance – the use of append helps us to reuse the first bag of t-inf.

Kiselyov et al. [3] has shown that a *MonadPlus* hides in implementations of logic programming system. Our BFS implementation is not an exception: append-map-inf/fair is like bind, but takes arguments in reversed order; none, unit, and append-inf/fair correspond to mzero, unit, and mplus respectively.

Functions implementing impure features are in Fig. 5. The first function, elim, takes a space s-inf and two continuations ks and kf. When s-inf contains some answers, ks is called with the first answer and the rest space. Otherwise, kf is called with no argument. Here 's' and 'f' means 'succeed' and 'fail' respectively. This

benchmark	size	DFS_i	DFS_{bi}	DFS_f	optimized BFS	Silvija's BFS
very-recursiveo	100000	579	793	2131	1438	3617
	200000	1283	1610	3602	2803	4212
	300000	2160	2836	-	6137	-
appendo	100	31	41	42	31	68
	200	224	222	221	226	218
	300	617	634	593	631	622
reverso	10	5	3	3	38	85
	20	107	98	51	4862	5844
	30	446	442	485	123288	132159
quine-1	1	71	44	69	-	-
	2	127	142	95	-	-
	3	114	114	93	-	-
quine-2	1	147	112	56	-	-
	2	161	123	101	-	-
	3	289	189	104	-	-
'(I love you)-1	99	56	15	22	74	165
	198	53	72	55	47	74
	297	72	90	44	181	365
'(I love you)-2	99	242	61	16	66	99
	198	445	110	60	42	64
	297	476	146	49	186	322

Table 1. The results of a quantitative evaluation: running times of benchmarks in milliseconds

function is like an eliminator of search space, hence the name. The remaining functions do the same thing as before.

6.2 comparison with the BFS of Seres et al. [5]

In this section, we compare the pure subset of our optimized BFS with the BFS found in Seres et al. [5]. We focus on the pure subset because Silvija's system is pure. Their system represents search spaces with streams of lists of answers, where each list is a bag.

To compare efficiency, we translate her Haskell code into Racket (See supplements for the translated code). The translation is direct due to the similarity in both logic programming systems and search space representations. The translated code is longer and slower than our BFS implementation. Details about difference in efficiency are in section 6.

We prove in Coq that the two BFSs are equivalent, i.e. (run n g) produces the same result (See supplements for the formal proof).

7 QUANTITATIVE EVALUATION

In this section, we compare the efficiency of search strategies. A concise description is in Table 1. A hyphen means running out of memory. The first two benchmarks are taken from Friedman et al. [2]. reverso is from Rozplokhas and Boulytchev [4]. Next two benchmarks about quine are modified from a similar test case in Byrd et al. [1]. The modifications are made to circumvent the need for symbolic constraints (e.g. ≠, absento). Our version generates de Bruijnized expressions and prevent closures getting into list. The two benchmarks differ in the $cond^e$ clause

order of their relational interpreters. The last two benchmarks are about synthesizing expressions that evaluate to '(I love you). This benchmark is also inspired by Byrd et al. [1]. Again, the sibling benchmarks differ in the $cond^e$ clause order of their relational interpreters. The first one has elimination rules (i.e. application, car, and cdr) at the end, while the other has them at the beginning. We conjecture that DFS_i would perform badly in the second case because elimination rules complicate the problem when running backward. The evaluation supports our conjecture.

In general, only DFS_i and DFS_{bi} constantly perform well. DFS_f is just as efficient in all benchmarks but very-recursiveo. Both BFS have obvious overhead in many cases. Among the three variants of DFS (they all have unfair conj), DFS_f is most resistant to clause permutation, followd by DFS_{bi} then DFS_i. Among the two implementation of BFS, ours constantly performs as well or better. Interestingly, every strategies with fair disj suffers in very-recursiveo and DFS_f performs well elsewhere. Therefore, this benchmark might be a special case. Fair conj imposes overhead constantly except in appendo. The reason might be that strategies with fair conj tend to keep more intermediate answers in the memory.

8 RELATED WORKS

Edward points out a disjunct complex would be 'fair' if it is a full and balanced tree Yang [6].

Silvija et al Seres et al. [5] also describe a breadth-first search strategy. We proof their BFS is equivalent to ours. But our code looks simpler and performs better in comparison with a straightforwrd translation of their Haskell code.

9 CONCLUSION

We analysis the definitions of fair disj and fair conj, then propose a new definition of fair conj. Our definition is orthogonal with completeness.

We devise two new search strategies (i.e. balanced interleaving DFS (DFS_{bi}) and fair DFS (DFS_f)) and devise a new implementation of BFS. These strategies have different features in fairness: $_{bi}$ has almost-fair disj and unfair conj. DFS_f has fair disj and unfair conj. BFS has both fair disj and fair conj.

Our quantitative evaluation shows that DFS_{bi} and DFS_f are competitive alternatives to DFS_i, the current search strategy, and that BFS is less practical.

We prove our BFS is equivalent to the BFS in Seres et al. [5]. Our code is shorter and runs faster than a direct translation of their Haskell code.

REFERENCES

[1] William E Byrd, Michael Ballantyne, Gregory Rosenblatt, and Matthew Might. 2017. A unified approach to solving seven programming problems (functional pearl). *Proceedings of the ACM on Programming Languages* 1, ICFP (2017).
[2] Daniel P. Friedman, William E. Byrd, Oleg Kiselyov, and Jason Hemann. 2018. *The Reasoned Schemer, Second Edition*.
[3] Oleg Kiselyov, Chung-chieh Shan, Daniel P Friedman, and Amr Sabry. 2005. Backtracking, interleaving, and terminating monad transformers:(functional pearl). *ACM SIGPLAN Notices* 40, 9 (2005), 192–203.
[4] Dmitri Rozplokhas and Dmitri Boulytchev. 2018. Improving Refutational Completeness of Relational Search via Divergence Test. In *Proceedings of the 20th International Symposium on Principles and Practice of Declarative Programming*. ACM, 18.
[5] Silvija Seres, J Michael Spivey, and C. A. R. Hoare. 1999. Algebra of Logic Programming.. In *ICLP*. 184–199.
[6] Edward Z. Yang. 2010. Adventures in Three Monads. *The Monad. Reader Issue 15* (2010), 11.