

# miniKanren with fair search strategies

KUANG-CHEN LU, Indiana University

WEIXI MA, Indiana University

DANIEL P. FRIEDMAN, Indiana University

The syntax of a programming language should reflect its semantics. When using a disjunction operator in relational programming, a programmer would expect all clauses of this disjunction to share the same chance of being explored, as these clauses are written in parallel. The existing multiarity disjunctive operator in miniKanren, however, prioritize its clauses by the order of which these clauses are written down. We have devised two new search strategies that allocate computational effort more fairly in all clauses.

## ACM Reference Format:

Kuang-Chen Lu, Weixi Ma, and Daniel P. Friedman. 2019. miniKanren with fair search strategies. 1, 1 (April 2019), 8 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 INTRODUCTION

When every sub-goal of a disjunction produces infinite states, the existing disjunctive operator allocates half computational effort to its first goal, quarter to the second, eighth to the third, and so on. The unfairness provides both opportunity and burden: miniKanren users can place more frequently used goal at the beginning to optimize their programs; however, it might be a catastrophe if a goal that generate many useless partial solution is placed before more important goals. Seasoned miniKanreners usually know how to utilize the unfairness to optimize their programs. However, we believe search strategies that is less sensitive to goal order can also be useful to little miniKanreners as well as seasoned ones. We propose two such search strategies, balanced interleaving DFS (biDFS) and breadth-first search (BFS), and observe how they affect the efficiency and the answer order of known miniKanren programs. The experiment is conducted with the miniKanren from *The Reasoned Schemer, 2nd Edition*.

## 2 RELATED WORKS

The unfairness of disjunction has been noticed by Seres et al. Their complete and fair search strategy is also named breadth-first search. Their BFS is fair and is similar to ours. However, their Haskell implementation cannot be translated to Scheme directly. This is partially due to the difference in calling conventions of host languages. Besides, their search space are infinite even when no answers exists. This is not feasible in miniKanren, where users can query all answers.

---

Authors' addresses: Kuang-Chen LuIndiana University; Weixi MaIndiana University; Daniel P. FriedmanIndiana University.

---

Unpublished working draft. Not for distribution.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).  
© 2019 Association for Computing Machinery.  
XXXX-XXXX/2019/4-ART \$15.00  
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

```

48 (defrel (repeato x out)
49   (conde
50     [(== '() out)]
51     [(fresh (res)
52       (== '(,x . ,res) out)
53       (repeato x res))]))
54
55
56

```

Fig. 1. repeato

### 3 FAIRNESS

A disjunction is fair if when a corresponding goal is queried, answers of lower costs come first. The *cost* of an answer is the number of relation applications needed to verify the answer.

Now we illustrate the costs of answers by running a miniKanren relation. Fig. 1 defines the relation **repeato** that relates a term *x* with a list whose elements are all *xs*.

Consider the following run of **repeato**.

```

65 > (run 4 q
66     (repeato '* q))
67 '(() (*) (* *) (* * *))

```

The above **run** generates 4 answers. All are lists of *\*s*. The order of the answers reflects the order miniKanren discovers them: the first answer in the list is first discovered. This order is not surprising: to generate the first answer, *'()*, miniKanren needs to apply **repeato** only once and the later answers need more recursive applications. In this example, the cost of each answer is the same as one more than the number of *\*s*: the cost of *'()* is 1, the cost of *'(\*)* is 2, and so on.

A query result (list of answers) is in the *cost-respecting* order if no answer occurs before another answer of a lower cost. In the above example, the result is cost-respecting. The iDFS search, however, does not generate cost-respecting answers in general. As an example, consider the following run of **repeato**.

```

77 > (run 12 q
78     (conde
79       [(repeato 'a q)]
80       [(repeato 'b q)]
81       [(repeato 'c q)]))
82 '(() (a) ()
83     (a a) () (a a a)
84     (b) (a a a a) (c)
85     (a a a a a) (b b) (a a a a a))
86

```

The results is not cost-respecting. For example, *'(a a)* occurs before *'(b)* while *'(a a)* is associated with a higher cost. iDFS strategy is the cause, since it prioritizes the first **conde** case considerably. When every **conde** case are equally productive, the iDFS strategy takes  $1/2^i$  answers from the *i*-th case, except the last case, which share the same portion as the second last one. In contrast, the same run with BFS produces answers in an expected order (Fig. ??).

```

92 > (run 12 q
93     (conde
94

```

```

95 (define (disj2 g1 g2)
96   (lambda (s)
97     (append-inf (g1 s) (g2 s))))
98
99 (define (append-inf s-inf t-inf)
100   (cond
101     ((null? s-inf) t-inf)
102     ((pair? s-inf)
103      (cons (car s-inf)
104            (append-inf (cdr s-inf) t-inf)))
105     (else (lambda ()
106              (append-inf t-inf (s-inf))))))
107
108
109

```

Fig. 2. disj2 and append-inf

```

111      [(repeato 'a q)]
112      [(repeato 'b q)]
113      [(repeato 'c q)])
114
115 '(() () ())
116 (a) (b) (c)
117 (a a) (b b) (c c)
118 (a a a) (b b b) (c c c))
119

```

#### 4 REPRESENTATION OF SEARCH SPACE

“Stream” is often used to name the representation of search space in miniKanren. However, the search space is more like a stream of list, because the information that a stream is thunk is employed to redirect search effort. If the search space is a normal stream, the information should have been used in a more trivial way.

In the rest of this paper, we call the `cars` of a stream its *mature part*, and the last `cdr` its *immature part*.

#### 5 WHY DISJ IS UNFAIR

The multiarity disjunction operator `disj` combine its sub-goals with `disj2`, right associatively. Interestingly, although `disj` is unfair, `disj2` is fair. The core functionality of `disj2` is completed by `append-inf` (Fig. reldisj2-and-append-inf). When both input streams are infinite, the resulting mature part contains only the mature part of `s-inf`. The whole `t-inf` goes to the resulting immature part. However, `t-inf` and `s-inf` are swapped in the delayed recursive call. Hence the search strategy spend computational effort evenly in two appended streams. As `disj` applies `disj2` right associatively, the left clauses are more frequently explored.

#### 6 BALANCED INTERLEAVING DFS

Our first solution, balanced interleaving DFS, is based on the observation that a disjunction can be viewed as a binary tree, where `disj2`s are nodes and sub-goals are leaves. In iDFS, the tree is in one of the most

```

142 (define (split ls k)
143   (cond
144     [(null? ls) (k '() '())]
145     [else (split (cdr ls)
146                  (lambda (l1 l2)
147                    (k l2 (cons (car ls) l1))))))])
148
149 (define (disj* gs)
150   (cond
151     [(null? gs) fail]
152     [(null? (cdr gs)) (car gs)]
153     [else
154      (split gs
155             (lambda (gs1 gs2)
156               (disj2 (disj* gs1)
157                      (disj* gs2))))]))
156
157 (define-syntax disj
158   (syntax-rules ()
159     [(disj g ...) (disj* (list g ...))]))
160
161

```

Fig. 3. balanced-disj

unbalanced forms, because `disj2` is applied right associatively. As `disj2` allocates computational effort interleavily to its two sub-goals, a disjunction allocates half computational effort to its first sub-goal, quarter to the second, eighth to the third, and so on.

The key idea of biDFS is to make disjunction trees balanced (Fig. 3). We change the `disj` macro, and introduce a function `disj*` and its helper `split`. `disj*` essentially constructs a balanced `disj2` tree. The `split` helper splits elements of `ls` into two lists of roughly the same length, then apply `k` to them.

## 7 BREADTH-FIRST SEARCH

In this section we change the search strategy to *breadth-first search* and optimize it. The whole process is completed in two steps. In the first step, from `mk-0` to `mk-1`, BFS is introduced. In the second step, `mk-1` to `mk-3`, BFS is optimized. The initial version, `mk-0`, is exactly the version in *The Reasoned Schemer, 2nd Edition*.

### 7.1 from `mk-0` to `mk-1`

In `mk-0` and `mk-1`, search spaces are represented by streams of answers. Streams can be finite or infinite. Finite streams are just lists. And infinite streams are improper lists, whose last `cdr` is a thunk returning another stream. We call the `cars` the *mature* part, and the last `cdr` the *immature* part.

Streams are cost respective when they are initially constructed by `==`. However, the `mk-0` version of `append-inf` (Fig. ??) breaks cost respectiveness if its first input stream, `s-inf`, is infinite. The resulting

```

189 (define (append-inf s-inf t-inf)
190   (cond
191     ((null? s-inf) t-inf)
192     ((pair? s-inf)
193      (cons (car s-inf)
194            (append-inf (cdr s-inf) t-inf)))
195     (else (lambda ()
196              (append-inf t-inf (s-inf))))))

```

Fig. 4. append-inf in mk-0

```

201 (define (append-inf s-inf t-inf)
202   (append-inf^ #t s-inf t-inf))
203
204 (define (append-inf^ s? s-inf t-inf)
205   (cond
206     ((pair? s-inf)
207      (cons (car s-inf)
208            (append-inf^ s? (cdr s-inf) t-inf)))
209     ((null? s-inf) t-inf)
210     (s? (append-inf^ #f t-inf s-inf))
211     (else (lambda ()
212              (append-inf (t-inf) (s-inf))))))

```

Fig. 5. append-inf in mk-1

mature part contains only the mature part of `s-inf`. The whole `t-inf` goes to the resulting immature part.

The `mk-1` version of `append-inf` (Fig. 5) restores cost-respectiveness by combining the mature parts in the fashion of `append`. This `append-inf` calls its helper immediately, with the first argument, `s?`, set to `#t`, which means `s-inf` in the helper is the `s-inf` in the driver. Two streams are swapped in the third `cond` clause, with `s?` flipped accordingly.

`mk-1` is not efficient in two aspects. `append-inf` need to copy all `cons` cells of two input streams when the first stream has a non-trivial immature part. Besides, `mk-1` computes answers of the same cost at once, even when only a portion is queried. We solve the two problems in the next subsections.

## 7.2 mk-3, optimized breadth-first search

We avoid generating same-cost answers at once by expressing BFS with a queue. The elements of the queue are delayed computation, represented by thunks. Every `mk-1` stream has zero or one thunk, so we have no interesting way to manage it. Therefore we change the representation of immature parts from thunks to lists of thunks. As a consequence, we also change the way to combine mature and immature part from `append` to `cons`.

After applying these two changes, stream representation becomes more complicated. It motivates us to set up an interface between stream functions and the rest of miniKanren. Listed in Fig. 6 are all functions being aware of the stream representation, but `take-inf` and its helper function, which is explained later. The first three functions are constructors: `empty-inf` constructs an empty stream; `unit-mature-inf` constructs a stream with one mature solution; `unit-immature-inf` constructs a stream with one thunk. The `append-inf` in `mk-3` is relatively straightforward compared with the `mk-1` version. `append-map-inf` is more tricky on how to construct the new immature part. We can follow the approach in `mk-0` and `mk-1` – create a new thunk which invoke `append-map-inf` recursively when forced. But then we need to be careful: if we construct the thunk when the old immature part is an empty list, the resulting stream might be infinitely unproductive. Besides, all solutions of the next lowest cost in `s-inf` must be computed when the thunk is invoked. However sometimes only a portion of these solutions is required to answer a query. To avoid the trouble and the advanced computation, we choose to create a new thunk for every existing thunk. The next four functions are used only by `ifte` and `once`. Uninterested readers might skip them. `null-inf?` checks whether a stream is exhausted. `mature-inf?` checks whether a stream has some mature solutions. `car-inf` takes the first solution out of a mature stream. `cdr-inf` drops the first solution of a mature stream. Finally, `force-inf` forces an immature stream to do more computation.

The last interesting function is `take-inf` (Fig. 7). The parameter `vs` is a list of solutions. The next two parameters, `P` and `Q`, together represent a queue. The first two `cond` lines are very similar to their counterparts in `mk-0` and `mk-1`. The third line runs when we exhaust all solutions. The fourth line re-shape the queue. The fifth and last line invoke the first thunk in the queue and use the mature part of the resulting stream, `s-inf`, as the new `vs`, and enqueueing `s-inf`'s thunks.

## 8 CONCLUSION

## ACKNOWLEDGMENTS

## REFERENCES

```

283 (define (empty-inf) '(() . ()))
284 (define (unit-mature-inf v) '((,v) . ()))
285 (define (unit-immature-inf th) '(() . (,th)))
286
287 (define (append-inf s-inf t-inf)
288   (cons (append (car s-inf) (car t-inf))
289         (append (cdr s-inf) (cdr t-inf))))
290
291 (define (append-map-inf g s-inf)
292   (foldr append-inf
293         (cons '()
294               (map (lambda (t)
295                     (lambda () (append-map-inf g (t))))
296                     (cdr s-inf)))
297         (map g (car s-inf))))
298
299 (define (null-inf? s-inf)
300   (and (null? (car s-inf))
301        (null? (cdr s-inf))))
302
303 (define (mature-inf? s-inf)
304   (pair? (car s-inf)))
305
306 (define (car-inf s-inf)
307   (car (car s-inf)))
308
309 (define (force-inf s-inf)
310   (let loop ((ths (cdr s-inf)))
311     (cond
312       ((null? ths) (empty-inf))
313       (else (let ((th (car ths)))
314                (append-inf (th)
315                             (loop (cdr ths))))))))
316
317
318
319
320
321
322
323
324
325
326
327
328
329

```

Fig. 6. Functions being aware of stream representation

```

330 (define (take-inf n s-inf)
331   (take-inf^ n (car s-inf) (cdr s-inf) '()))
332
333 (define (take-inf^ n vs P Q)
334   (cond
335     ((and n (zero? n)) '())
336     ((pair? vs)
337      (cons (car vs)
338            (take-inf^ (and n (sub1 n)) (cdr vs) P Q)))
339     ((and (null? P) (null? Q)) '())
340     ((null? P) (take-inf^ n vs (reverse Q) '()))
341     (else (let ([th (car P)])
342              (let ([s-inf (th)])
343                (take-inf^ n (car s-inf)
344                             (cdr P)
345                             (append (reverse (cdr s-inf)) Q)))))))
346
347
348

```

Fig. 7. take-inf in mk-3-1